

MODEL DISAGREEMENT PROTOCOLS: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study ensembles and world models for robots. The motivating bet is: Define action protocols that distinguish useful model disagreement from noise.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For ensembles and world models for robots, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2026a;b;c;d; 2016; 2026e). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.545	0.455
observed_only_wam	hidden_mode_shift	0.133	0.867
augmentation_as_data_wam	nominal	0.615	0.385
augmentation_as_data_wam	hidden_mode_shift	0.287	0.713
scalar_uncertainty_wam	nominal	0.680	0.320
scalar_uncertainty_wam	hidden_mode_shift	0.360	0.640
proposed_branch_mechanism	nominal	0.777	0.223
proposed_branch_mechanism	hidden_mode_shift	0.537	0.463

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Model Disagreement Protocols argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Agile science: creating useful products for behavior change in the real world, 2016. <https://doi.org/10.1007/s13142-016-0395-7>.
- Maniparena: Comprehensive real-world evaluation of reasoning-oriented generalist robot manipulation, 2026a. <http://arxiv.org/abs/2603.28545v1>.
- Benchmarking robust machine learning models under data imperfections in real-world data science scenarios, 2026b. <https://doi.org/10.61453/jods.v20260103>.
- Failure-aware rl: Reliable offline-to-online reinforcement learning with self-recovery for real-world manipulation, 2026c. <http://arxiv.org/abs/2601.07821v1>.
- Evaluating large language model ensembles as probabilistic forecasters: An empirical study on live real-world events, 2026d. <https://doi.org/10.2139/ssrn.6709818>.
- Worldbench: Disambiguating physics for diagnostic evaluation of world models, 2026e. <http://arxiv.org/abs/2601.21282v1>.